
COMP604 – Operating Systems

Unofficial course manual

Written by Zach Barrett

Table of contents

| | |
|---|----------|
| 1. Introduction | 3 |
| 1.1. Additional resources | 3 |
| 2. Class Overview | 4 |
| 2.1. Assessment Overview | 4 |
| 2.1.1. Assignment One | 4 |
| 2.1.2. Assignment Two | 4 |
| 2.1.3. Assignment Three | 4 |
| 2.2. Software tools | 4 |
| 2.2.1. Mac setup | 5 |
| 2.2.2. Ubuntu setup | 5 |
| 2.2.3. Windows setup | 6 |
| 3. What you need to know about C programming | 7 |
| 4. XV6 for Dummies | 8 |
| 4.1. General structure | 8 |
| 4.2. System Calls | 8 |
| 4.2.1. How to add a system call | 8 |
| 5. Assignment Notes | 9 |
| 5.1. Assignment One | 9 |
| 5.1.1. Question One | 9 |
| 5.1.2. Question Two | 9 |
| 5.1.3. Question Three | 10 |
| 5.2. Assignment Two | 11 |
| 5.2.1. Question One | 11 |
| 5.2.2. Question Two | 11 |
| 5.2.3. Question Three | 11 |
| 5.2.4. Question Four | 11 |
| 5.2.5. Question Five | 11 |
| 5.3. Assignment Three | 11 |
| 5.3.1. Question One | 11 |
| 5.3.2. Question Two | 11 |
| 5.3.3. Question Three | 11 |
| 5.3.4. Question Four | 11 |

1. Introduction

This document is a collection of my notes and knowledge related to the class “Operating Systems” at AUT, written for someone taking this class. While largely intended as a way of writing down everything I know for my own benefit, my hope is that it will be of use for you as well.

1.1. Additional resources

This class recommends the use of the OSTEP (Operating Systems, Three Easy Pieces) textbook. It can be found here: [OSTEP](#) This is an excellent resource, and I highly recommend using this.

2. Class Overview

2.1. Assessment Overview

There are three assignments in this class, with assignment two being the longest and most difficult.

| Assessment Item | Weighting |
|------------------|-----------|
| Assignment One | 30% |
| Assignment Two | 40% |
| Assignment Three | 30% |

2.1.1. Assignment One

This assignment covered shell scripting, C programming, and a simple xv6 system call. Easiest by far, do not use your extension here. When I took it, this assignment had 3 questions.

2.1.2. Assignment Two

This assignment covered scheduling algorithms, processes, and more advanced system calls. This is the longest assignment, using your extension here is fine. When I took it, this assignment had 5 questions. All programming questions were done in xv6.

2.1.3. Assignment Three

This assignment covered semaphores, inodes, RAID, storage in xv6. This assignment was harder than assignment one but easier than assignment two. If you can save your extension for this assignment then great, but dont feel bad about spending it on assignment two.

2.2. Software tools

The later portion of this course makes use of the XV6 operating system. Because XV6 runs on a different architecture than most computers, it can only be run using a virtual machine. Thankfully, a free and easy to use virtual machine exists called QEMU. The next sections will show you how to install this software stack.

If you are able to get access to a linux or mac computer, doing so will make your life much easier as you will be able to run QEMU directly and not through another virtual machine (Unlike windows, where you will need to use WSL).

If you have to use WSL, I highly recommend installing Visual Studio Code and the WSL extension for it, allowing you to edit the xv6 files directly through the editor instead of having to use either copy them in/out or having to use a command line editor like nano or vim. (Although I also recommend you learn how to use vim keybindings.)

This information is a combination of the “Software setup” section on canvas and my own experiences.

2.2.1. Mac setup

1. To compile the software tools necessary to run XV6, you must first install xcode.

```
xcode-select --install
```

This will take some time.

2. Next, ensure the Homebrew package manager is installed. If it is not, install with the following command:

```
/bin/bash -c "$(curl -fsSL https://raw.githubusercontent.com/Homebrew/install/HEAD/install.sh)"
```

(If you don't already use homebrew, you should. It's extremely useful.)

3. Then, use Homebrew to install the risc-v compiler toolchain:

```
brew tap riscv/riscv  
brew install riscv-tools
```

With this step, sometimes people (me, I had this issue) have issues with compiling the riscv toolchain. If that occurs try changing your OSX version to OSX Sonoma, as there is a precompiled binary available for that version. Once that's done try running those commands again, Homebrew will handle it for you.

4. Next you must add the toolchain directory to your "path", which is where your computer looks for programs if they are not in the current directory. This can be done using any text editor. To do so, you must open a file called .zshrc (This is the configuration file for the default shell on mac, if you have changed the shell I assume you know which **other** file to change instead.)

Then append one of the following lines to the bottom of the file depending on your situation

```
# Use the below line on M series macs  
export PATH=$PATH:/usr/local/opt/homebrew/Cellar/riscv-gnu-toolchain/master/bin  
  
# Use this one for intel series macs  
export PATH=$PATH:/usr/local/Cellar/riscv-gnu-toolchain/master/bin  
  
# If you are on an M series mac and you get an error message at the end of  
# setup saying the computer can't find the compiler toolchain or something,  
# replace the line you added with this one:  
export PATH=$PATH:/opt/homebrew/Cellar/riscv-gnu-toolchain/main.reinstall/bin
```

5. Finally, install QEMU.

```
brew install qemu
```

If you have issues with this step ensure that the path was set correctly.

2.2.2. Ubuntu setup

Getting this running is very easy on ubuntu. Note that the university computers with WSL enabled also have ubuntu installed.

To setup QEMU on ubuntu, simply run the following commands:

```
sudo apt update && sudo apt upgrade  
sudo apt install build-essential gdb-multiarch qemu-system-misc gcc-riscv64-linux-gnu binutils-riscv64-linux-gnu
```

If that throws an error, try separating all of the packages into separate lines in a text file (Let's call it packs.txt), and running the following command:

```
cat packs.txt | xargs sudo apt install
```

(Learn about the xargs command by the way, very useful!)

2.2.3. Windows setup

You will need to use WSL (Windows Subsystem for Linux) as windows cannot run QEMU directly.

1. Start by opening a command prompt and running the following command:

```
wsl --install -d ubuntu-22.04
```

This will install WSL and ubuntu.

2. Reboot your computer, open bios and ensure that hardware virtualisation is enabled. (look up how to do this, it varies by motherboard manufacturer)
3. Once your computer has rebooted, open a command prompt and run:

```
ubuntu
```

This will open an ubuntu command prompt. From here setup proceeds the same as in the ubuntu setup section.

Once you have completed that, I recommend setting up an editor capable of opening and editing files in WSL. This can be done using Visual Studio Code and an extension. [Here is the link to that extension](#)

3. What you need to know about C programming

4. XV6 for Dummies

4.1. General structure

4.2. System Calls

4.2.1. How to add a system call

5. Assignment Notes

5.1. Assignment One

5.1.1. Question One

This question required me to write a bash script which took two arguments, a filename and the path to a directory. Then, it looked for all all lines in that file which contained instances of the word “special”, and appended those lines to an output file called special.txt in that directory.

It also had to do the following:

1. Display an error message if it wasn't passed exactly two arguments.
2. If the specified file didn't exist, it also had to display an error message.
3. If the specified output directory didn't exist, it must create it.
4. If a file called “special.txt” already exists, the lines must be appended to it.

This was reasonably easy, finding lines which contain a string can be done with the `grep` command line utility.

5.1.2. Question Two

This question was substantially harder. I needed to implement a linked list in C, which will be used to store “process control blocks”, which store the name and PID of a process.

First I needed to create a structure to store the PCB:

```
typedef struct PCB {
    int PID;    // Integer for pid
    char* name; // String for name
}
```

Then, I create a structure for the list item:

```
typedef struct List_item {
    struct List_item* next; // This stores the address of next item in the list
    struct PCB* referredPCB; // This stores the address of the PCB associated with this list node
} List_item
```

First was the insert function:

```
int insert(struct List_item *listhead, struct List_item *insertItem)
```

This function took the listhead and a list item, and inserted it into the correct position depending on the PCB's PID.

Next was the delete function:

```
int delete(struct List_item *listhead, int pid)
```

Which unlinks the node with the specified PID. One thing worth noting is that `delete()` is a reserved keyword in C++, not C. But because the commonly used clang code checker works for both C and C++, some editors will throw an error if you try to make a function called `delete` in C. Just ignore it, it will compile just fine.

Last was the list printing function:

```
printList(List_item* listhead)
```

This function walks the linked list and prints all of the list items out.

Additionally, the functionality of this program had to be split across three different files. These were `main.c`, `listfunc.c`, and a header file (Which I called `listfunc.h`)

The main hard part of this question is figuring out a way to iterate over the list. Turns out the best way to do this is to create a `List_item` pointer, and then setting it equal to the address of the head node. Then, to move to the next node you simply set it equal to the referenced node's next node.

Something like this:

```
List_item listhead;
// Add a bunch of child nodes here.
List_item* cursor = &listhead;

while(cursor->next != NULL) { // Make sure you dont try to do a null pointer deref!
    cursor = cursor->next;

    // You can do whatever you want to the node here,
    // I am printing its name as an example.
    printf("%s", cursor->PCB->name)
}
```

5.1.3. Question Three

This question involved creating a new system call for `xv6`, and a user process. This system call needed to print a list of active processes, and print information about them. Thankfully, there is a function in `proc.c` called `procdump()` which does almost all of what we need it to, we just need to add something that prints the process' memory size.

The user program should also take an option `-r`, which if used should only print this information about a process if it is actively running (`p->state == RUNNING.`)

There is a process pointer called 'p', and if you get the value of `p->sz` we can get this information.

From this point I was able to create the user program, which took the input from the user, determined if the `-r` option was called, and invoked the `psc` system call with that option. While this approach was effective, I recommended that my friends simply made two sytem calls (one for if you are printing all processes and another for if you were just printing the running ones) because it was easier.

Other than that, the system call was implemented as normal, but passing the `-r` option required additonal work.

5.2. Assignment Two

5.2.1. Question One

5.2.2. Question Two

5.2.3. Question Three

5.2.4. Question Four

5.2.5. Question Five

5.3. Assignment Three

5.3.1. Question One

5.3.2. Question Two

5.3.3. Question Three

5.3.4. Question Four